

# Hadoop Map Reduce and YARN

## Lecture 4

Sharun Akter khushbu

Lecturer, Dept.of CSE

Daffodil International University

# Agenda for today's Session

edure

1. What is Hadoop MapReduce?
2. MapReduce In Nutshell
3. Advantages of MapReduce
4. Hadoop MapReduce Approach with an Example
5. Hadoop MapReduce/YARN Components
6. YARN With MapReduce
7. Yarn Application Workflow
8. MapReduce Program with Hands On



## 2 main Hadoop Components

*Storage*

*Processing*



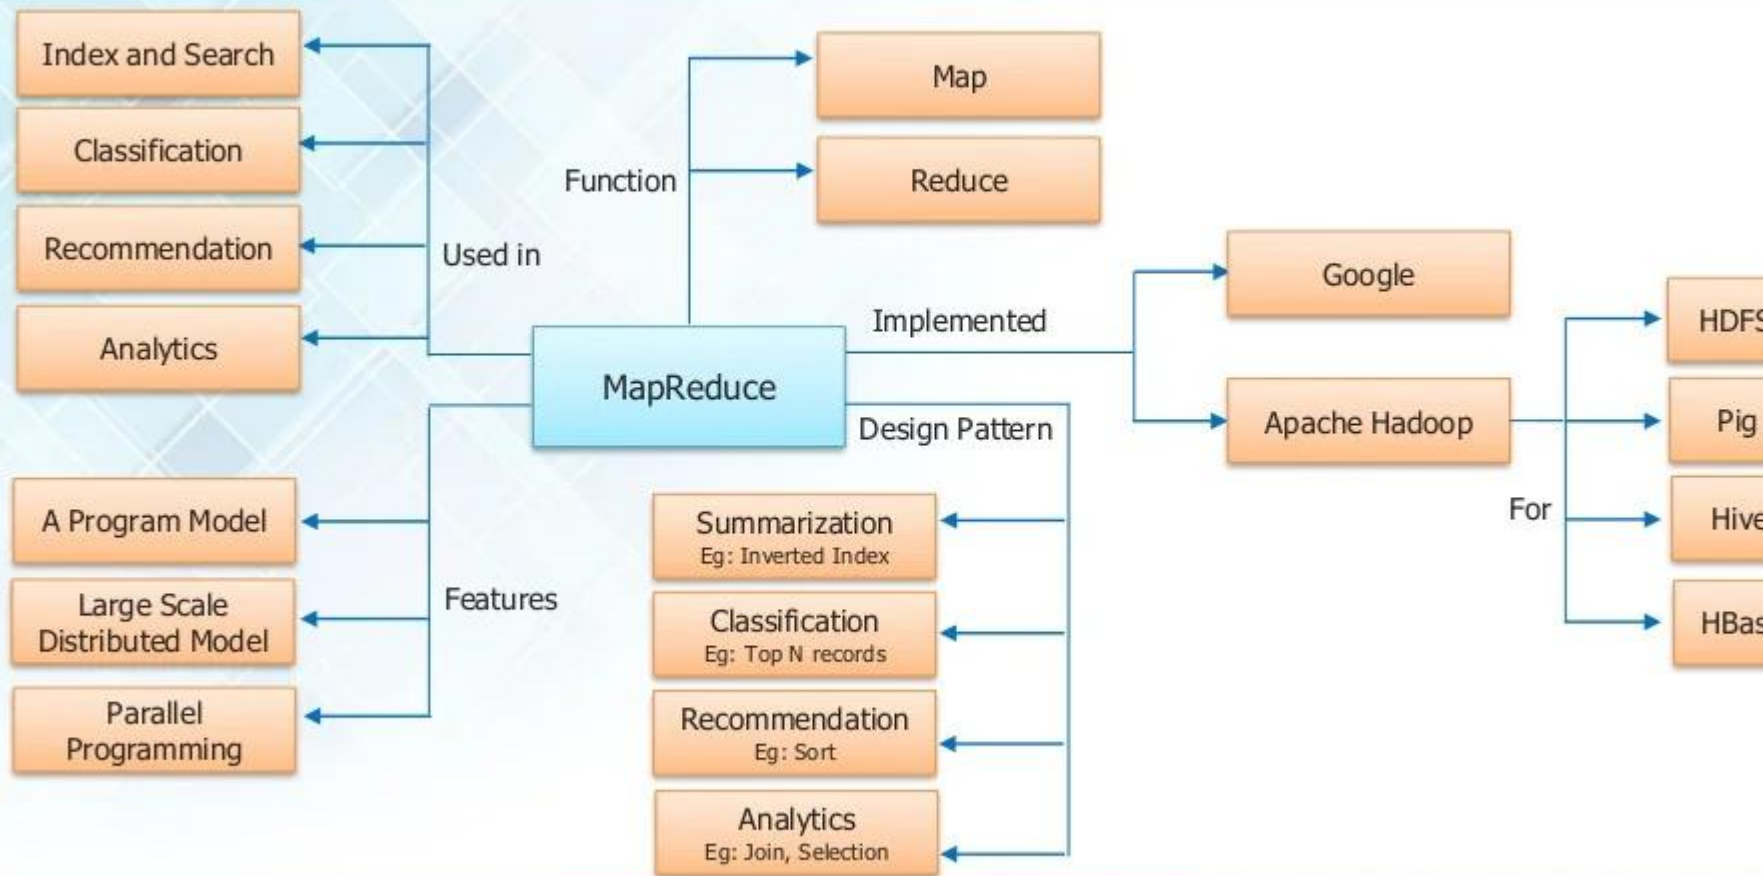
# MapReduce: Data Processing Using Programming

edure



- *Hadoop MapReduce is the processing component of Apache Hadoop*
- *It processes data parallelly in distributed environment*





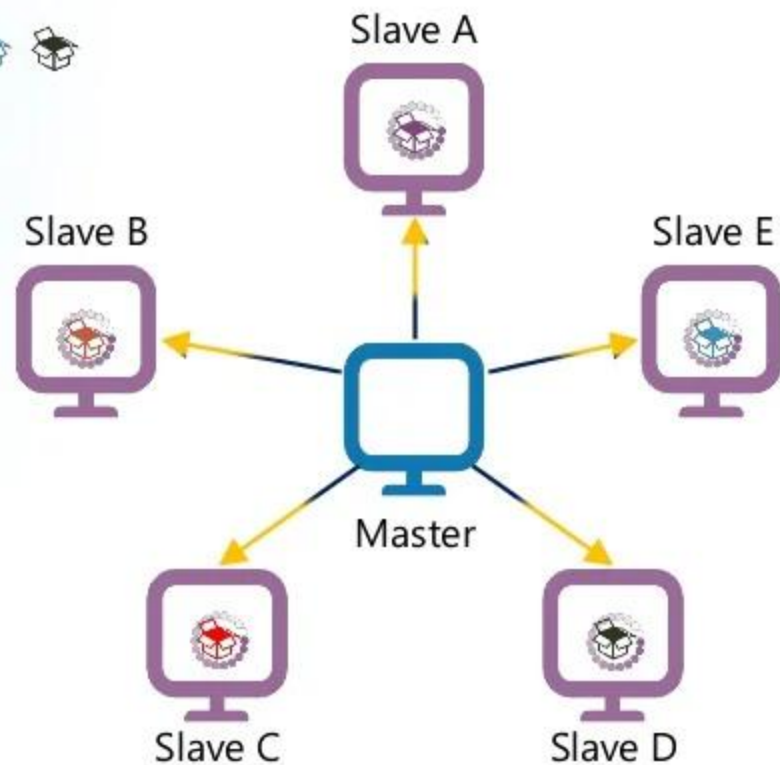
## 2 Biggest Advantages of MapReduce



# Advantage 1: Parallel Processing

Data → 

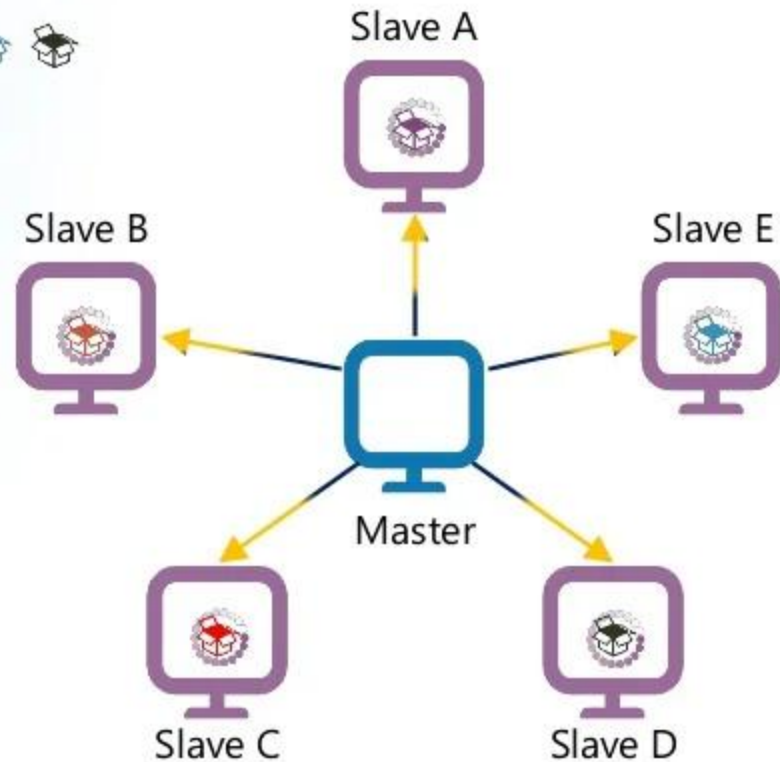
- Data is processed in parallel
- Processing becomes fast



## Advantage 2: Data Locality - Processing to Storage edure

Data → 

- Moving Data to processing is very costly
- In MapReduce, we move processing to Data



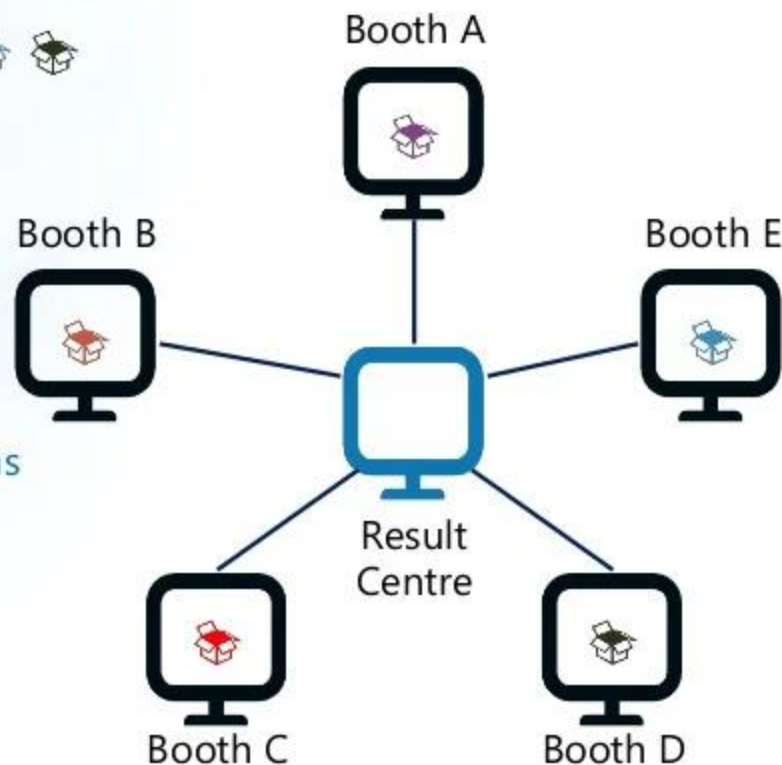


## Traditional vs MapReduce Way

Data → 

## Election Votes Casting


- Votes is stored at different Booths
- Result Centre has the details of all the Booths



Data → 

## Counting – Traditional Approach

- Votes are moved to Result Centre for counting
- Moving all the votes to Centre is costly
- Result Centre is over-burdened
- Counting takes time




# Hadoop MapReduce To the Rescue!

edure

Data → 

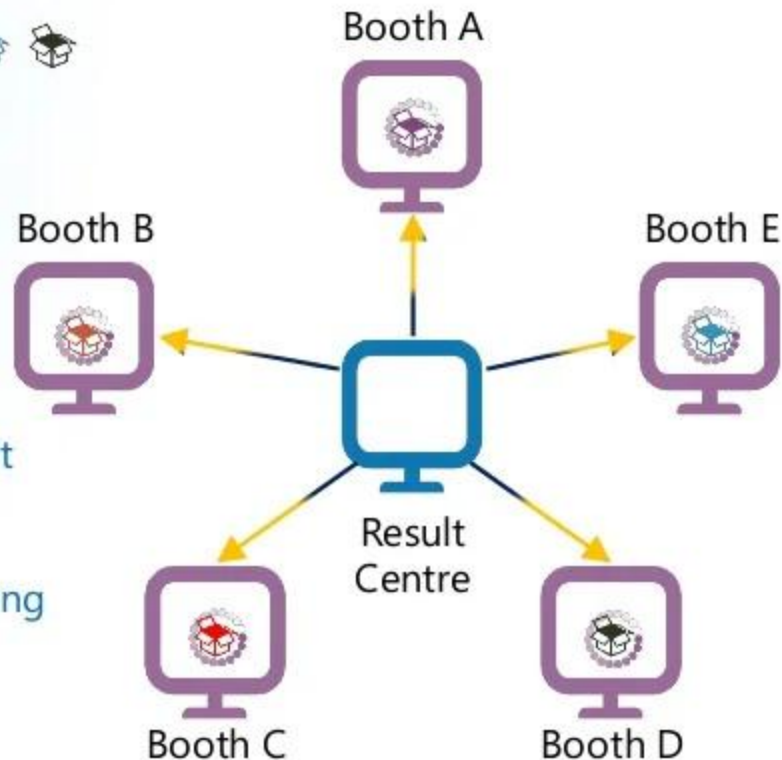
**✗ Hadoop MapReduce Doesn't Follow This Approach**



Votes → 

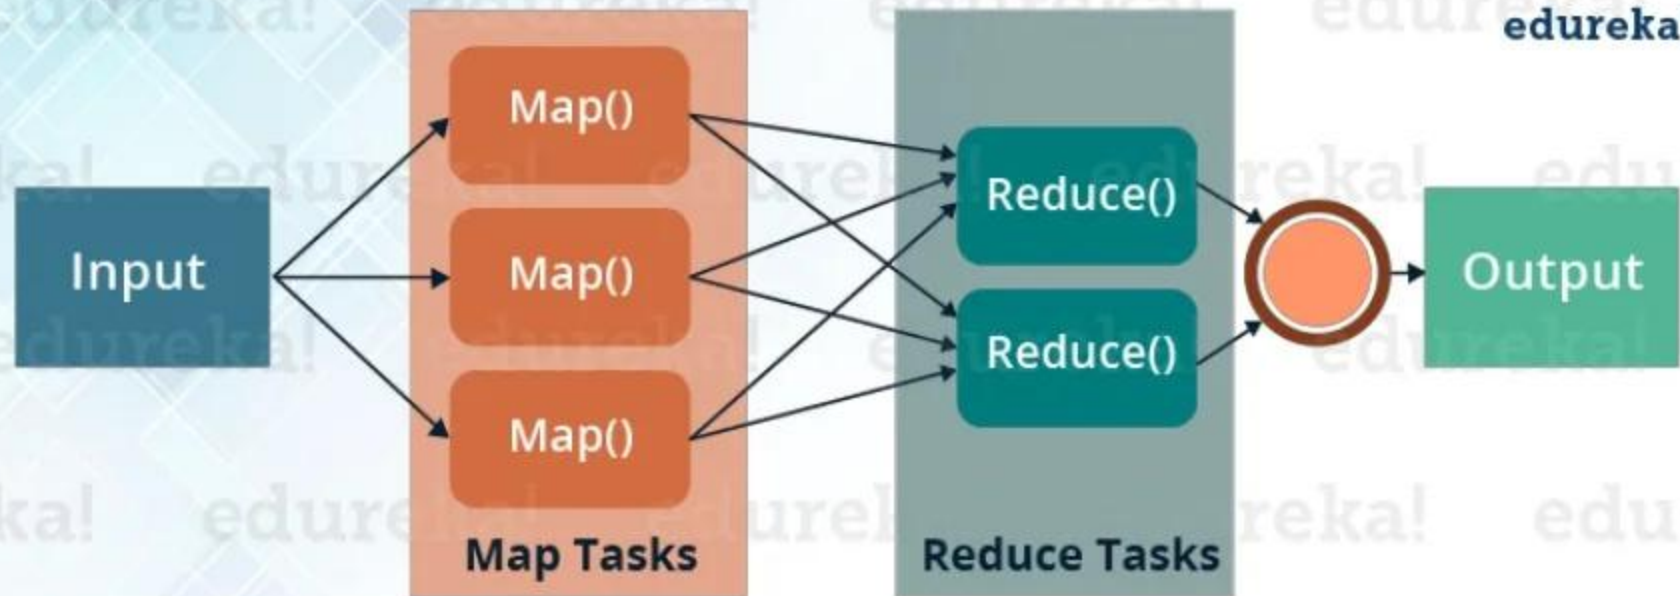
## Counting – MapReduce Approach

- Votes are counted at individual booths
- Booth-wise results are sent back to the result centre
- Final Result is declared easily and quickly using this way




# MapReduce In Detail





# Anatomy of a MapReduce Program

edure



Map:

(K1, V1)

List (K2, V2)

Reduce:


(K2, list (V2))

List (K3, V3)

Let us take an example to understand  
MapReduce Way

## The Overall MapReduce Word Count Process

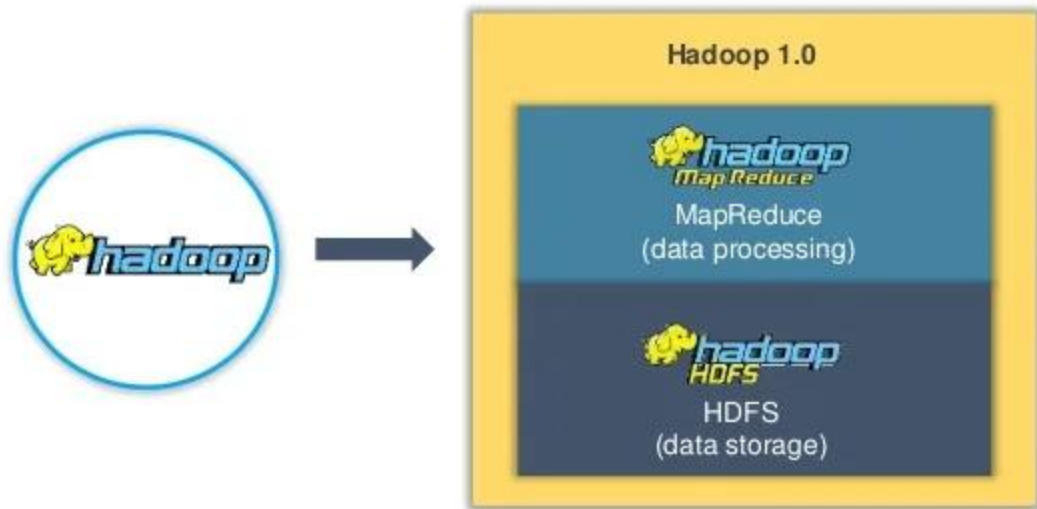
edureka!



Hadoop 1.0 (MR 1)




# Hadoop 1.0 (MR 1)



In Hadoop 1.0, **MapReduce** performed both **data processing** and **resource management**



Data processing



Resource management



# Hadoop 1.0 (MR 1)



MapReduce consisted of  
Job Tracker and Task Tracker



Job  
Tracker

Allocated resources, performed  
scheduling and monitored jobs

Assigned map and reduce tasks to jobs  
running on Task Trackers

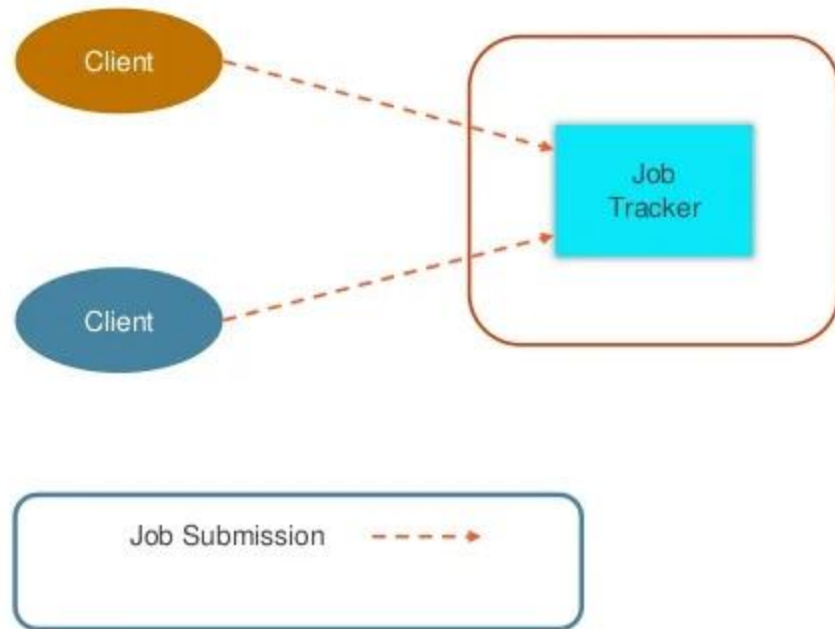
Task  
Tracker

Task Trackers processed the jobs

Task Trackers reported their progress  
to the Job Tracker

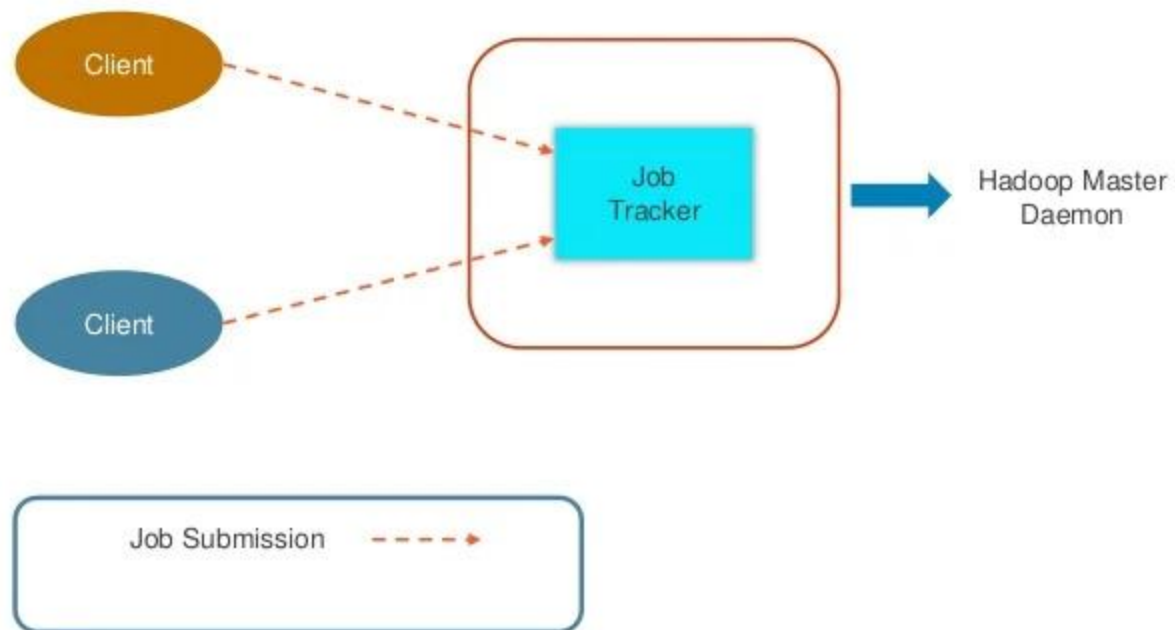
# Hadoop 1.0 (MR 1)

---







# Hadoop 1.0 (MR 1)

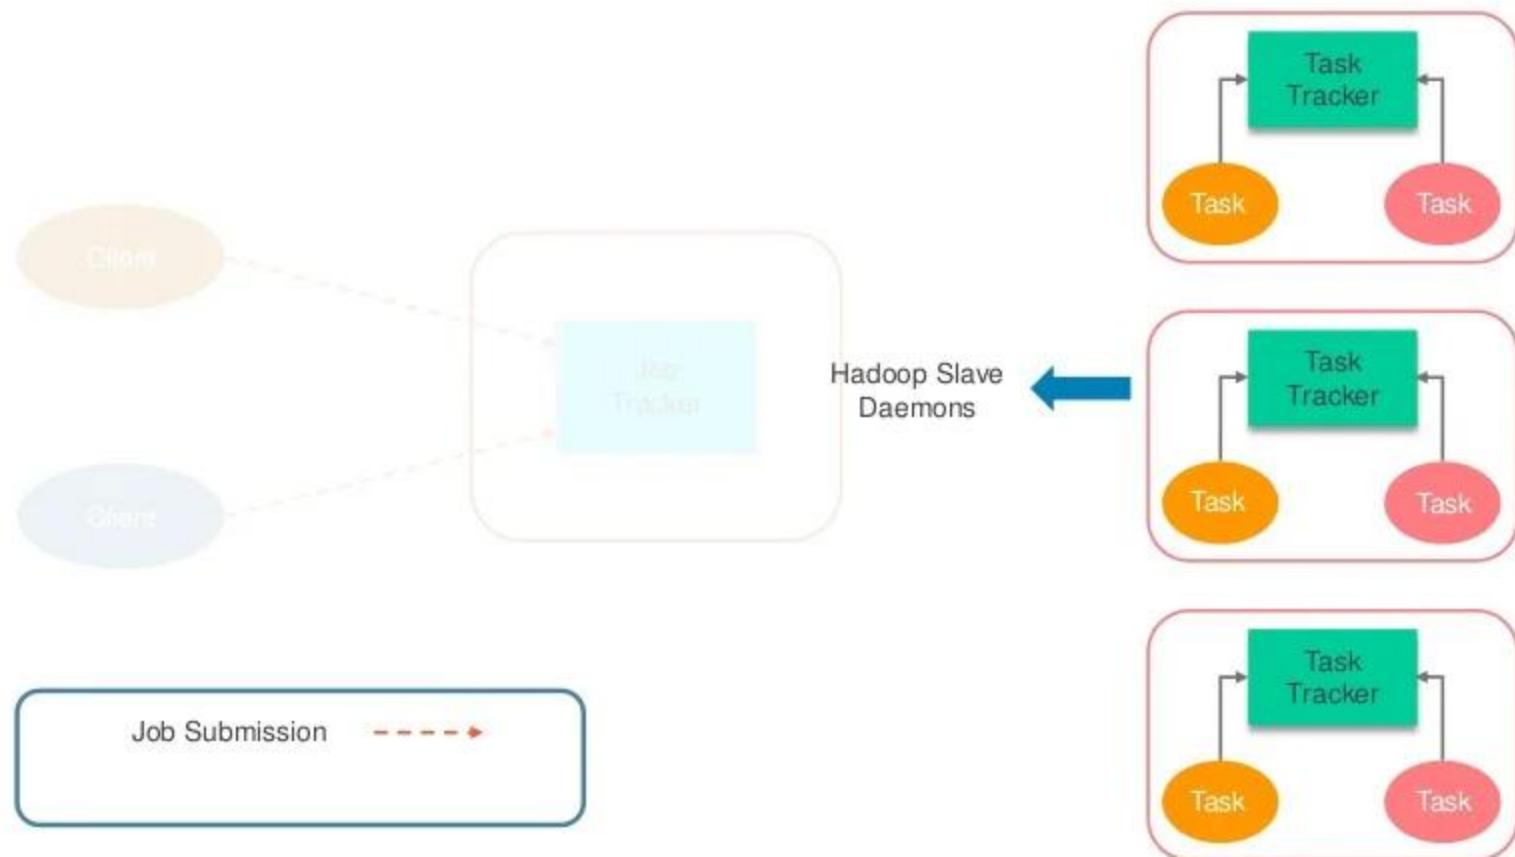
---




# Hadoop 1.0 (MR 1)



# Hadoop 1.0 (MR 1)

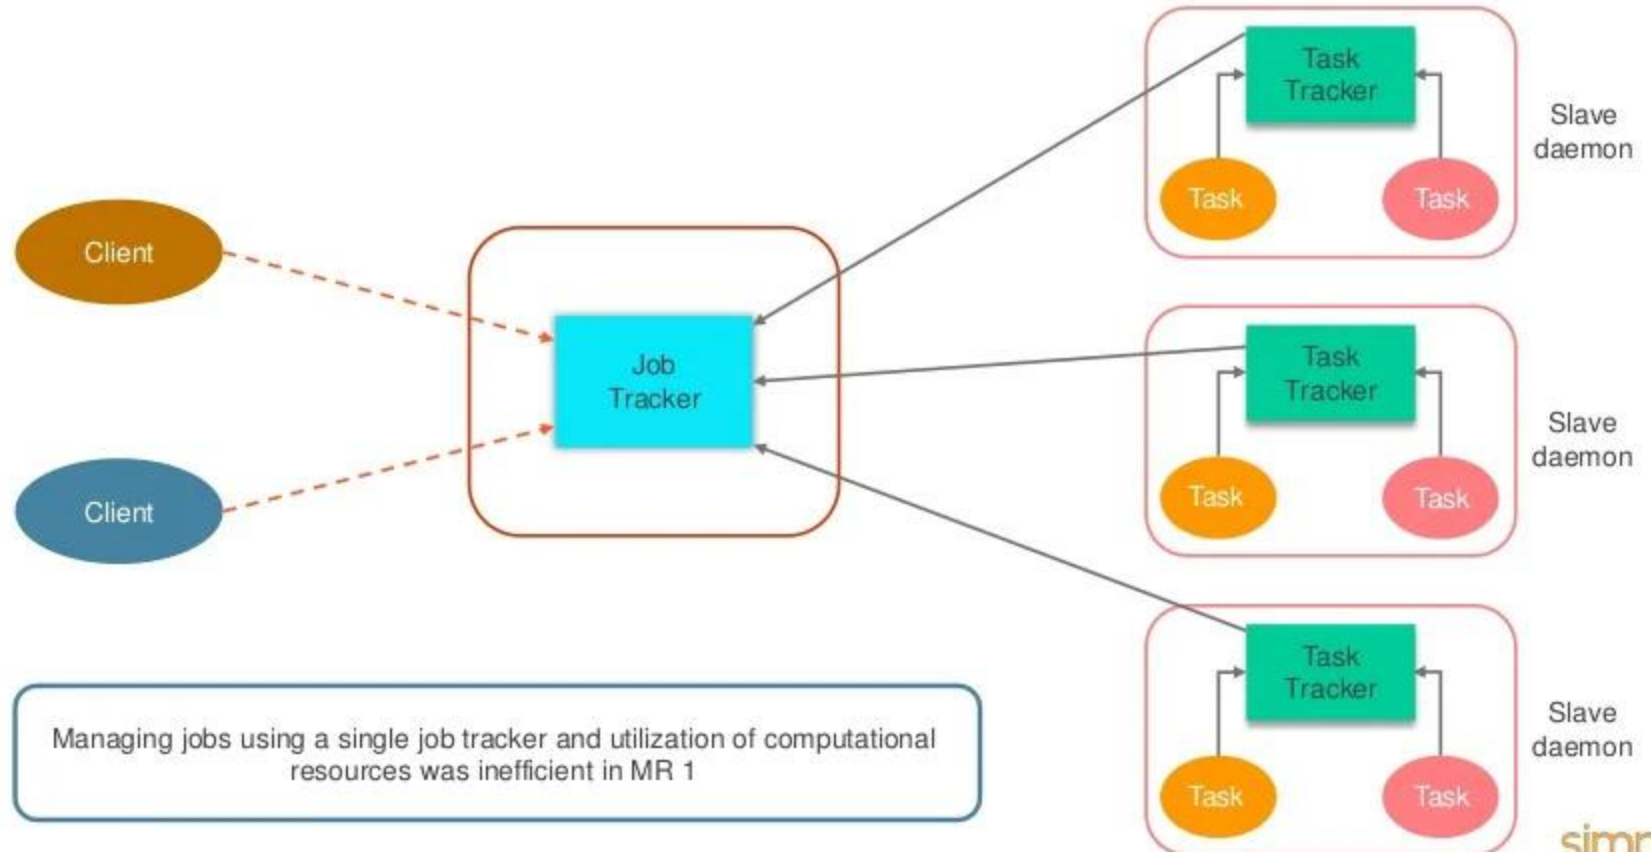


# Hadoop 1.0 (MR 1)





# Hadoop 1.0 (MR 1)



# Limitations of Hadoop 1.0 (MR 1)

1

## Scalability

Due to a [single JobTracker](#), scalability became a bottleneck.

Cannot have a cluster size of more than [4000 nodes](#) and cannot run more than [40000 concurrent tasks](#)



# Limitations of Hadoop 1.0 (MR 1)

1

Scalability

Due to a single JobTracker, scalability became a bottleneck.

Maximum cluster size - 4000 nodes  
Maximum concurrent tasks - 40000

2

Availability issue

JobTracker is **single point of failure**. Any failure kills all queued and running jobs. Jobs need to be resubmitted by users

# Limitations of Hadoop 1.0 (MR 1)

3

Resource Utilization

Due to predefined number of **map**  
and **reduce slots** for each  
TaskTracker, **resource utilization**  
issues occur



# Limitations of Hadoop 1.0 (MR 1)

3

Resource Utilization

Due to pre-defined number of map and reduce slots for each TaskTracker, resource utilization issues occur.

4

Limitations in running non-MapReduce applications

Problem in performing **real-time analysis** and running **Ad-hoc query** as MapReduce is batch driven

# MapReduce Using Yarn




Need for YARN



# Need for YARN


## Before YARN

Hadoop 1.0

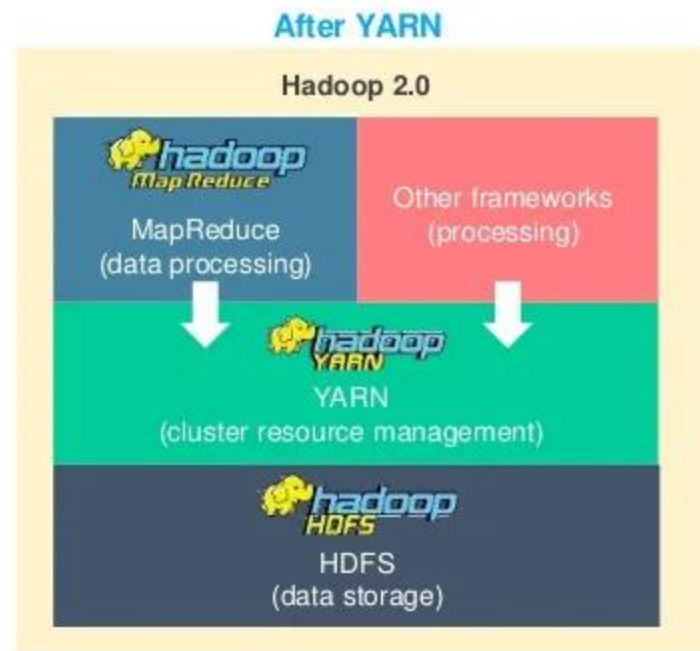


Designed to run MapReduce jobs only and had issues in scalability, resource utilization, etc.

# Need for YARN



Designed to run MapReduce jobs only and had issues in scalability, resource utilization, etc.



YARN solved those issues and users could work on multiple processing models along with MapReduce

## Solution - Hadoop 2.0 (YARN)



Scalability



Can have a cluster size of more than 10,000 nodes and can run more than 1,00,000 concurrent tasks

## Solution - Hadoop 2.0 (YARN)



### Scalability



Can have a cluster size of more than 10,000 nodes and can run more than 1,00,000 concurrent tasks

### Compatibility



Applications developed for Hadoop 1 runs on YARN without any disruption or availability issues

# Solution - Hadoop 2.0 (YARN)



## Scalability




Can have a cluster size of more than 10,000 nodes and can run more than 1,00,000 concurrent tasks

## Compatibility



Applications developed for Hadoop 1 runs on YARN without any disruption or availability issues

## Resource utilization



Allows dynamic allocation of cluster resources to improve resource utilization

# Solution - Hadoop 2.0 (YARN)



## Scalability




Can have a cluster size of more than 10,000 nodes and can run more than 1,00,000 concurrent tasks

## Compatibility



Applications developed for Hadoop 1 runs on YARN without any disruption or availability issues

## Resource utilization



Allows dynamic allocation of cluster resources to improve resource utilization

## Multitenancy



Can use open-source and propriety data access engines and perform real-time analysis and running ad-hoc query





*Applications Run Natively **IN** Hadoop*

**BATCH**  
(MapReduce)

**INTERACTIVE**  
(Text)

**ONLINE**  
(HBase)

**STREAMING**  
(Storm, S4, ...)

**GRAPH**  
(Giraph)

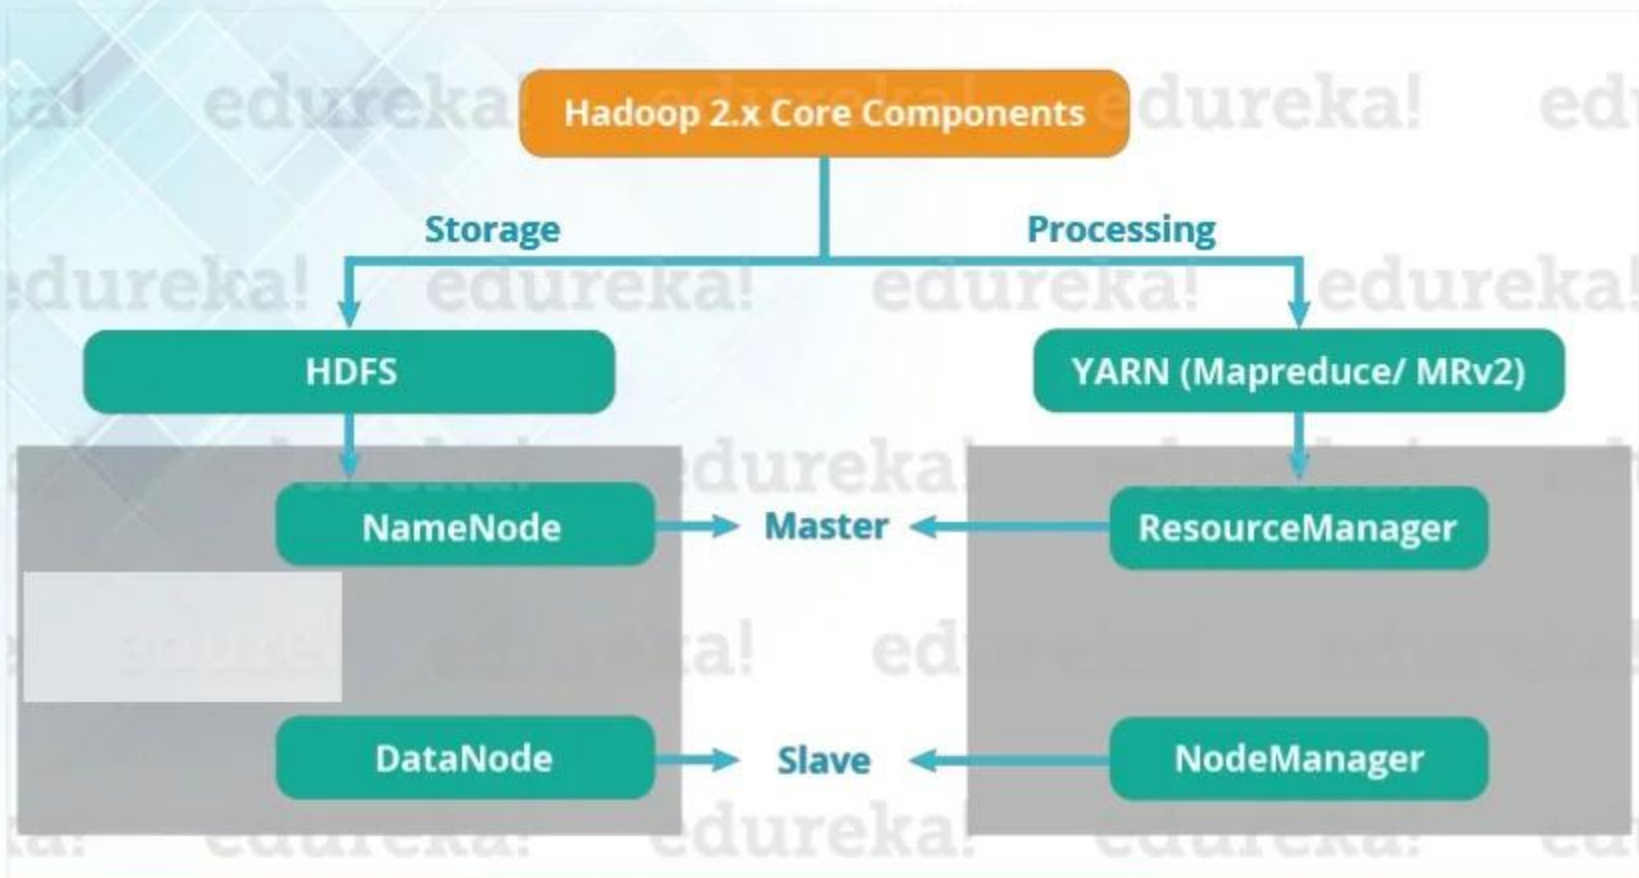
**IN-MEMORY**  
(Spark)

**HPC MPI**  
(OpenMPI)

**OTHER**  
(Search)  
(Weave..)

**YARN** (Cluster Resource Management)

**HDFS2** (Redundant, Reliable Storage)



## → Client

- » Submits a MapReduce Job

## → Resource Manager

- » Cluster Level resource manager
- » Long Life, High Quality Hardware

## → Node Manager

- » One per Data Node
- » Monitors resources on Data Node

## → Job History Server

- » Maintains information about submitted MapReduce jobs after their ApplicationMaster terminates

## → ApplicationMaster

- » One per application
- » Short life
- » Coordinates and Manages MapReduce Jobs
- » Negotiates with Resource Manager to schedule tasks
- » The tasks are started by NodeManager(s)

## → Container

- » Created by NM when requested
- » Allocates certain amount of resources (memory, CPU etc.) on a slave node

# YARN Architecture




# YARN Architecture

---







# YARN Architecture

---




# YARN Architecture







# YARN Architecture



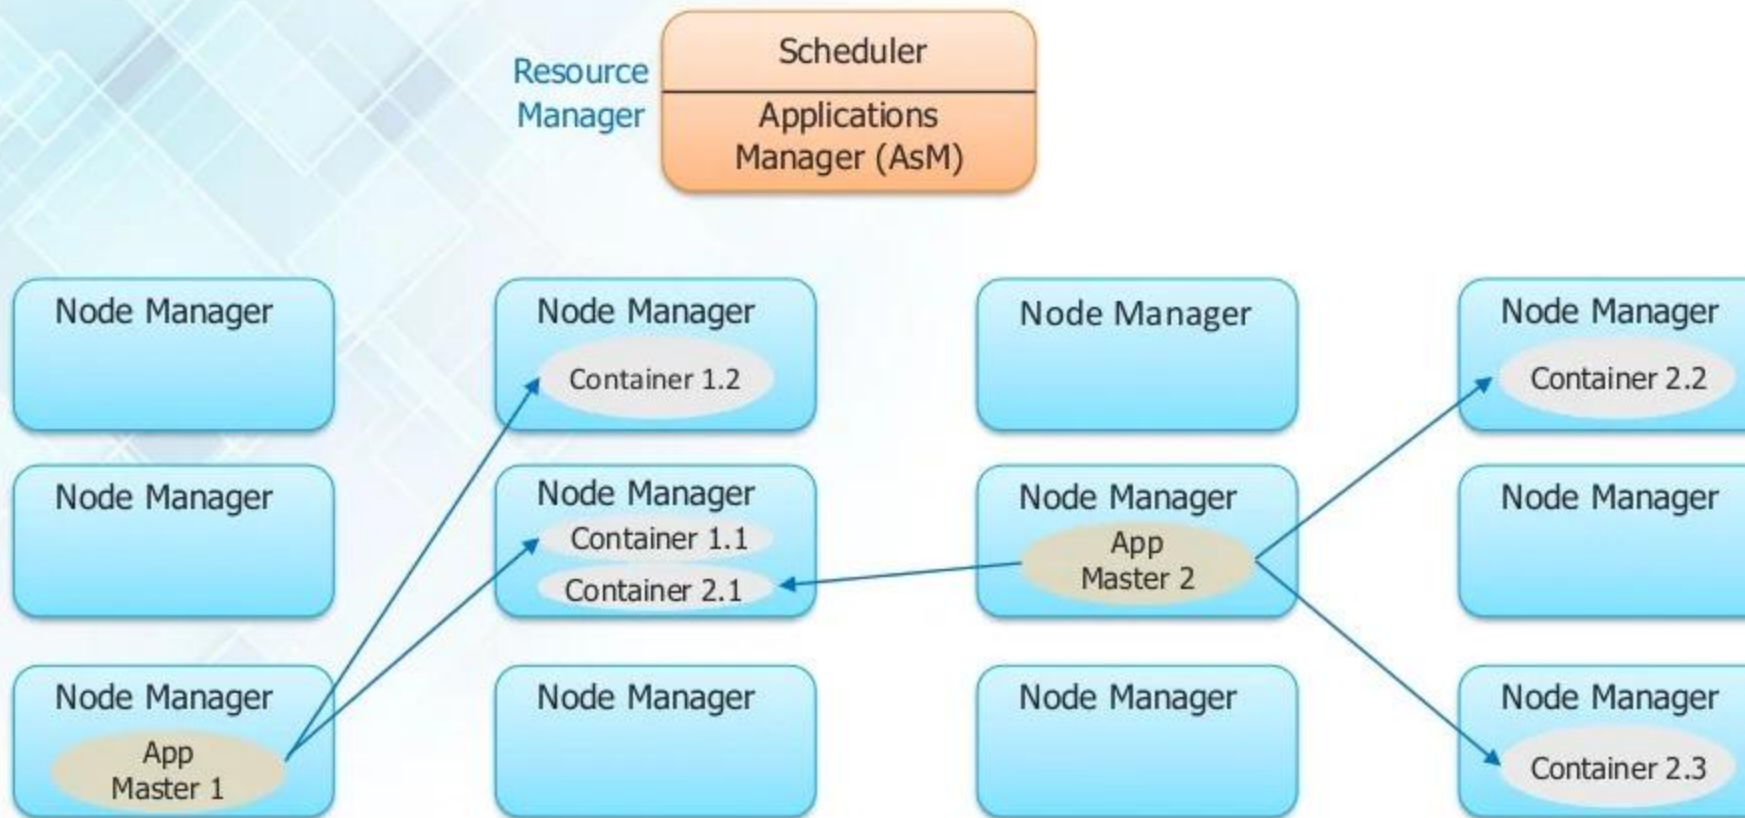
# YARN Architecture



# YARN Architecture




# YARN Application Workflow in MapReduce




→ Execution Sequence :

1. Client submits an application



→ Execution Sequence :


1. Client submits an application
2. RM allocates a container to start AM





→ Execution Sequence :


1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM



# Application Workflow

→ Execution Sequence :


1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM
4. AM asks containers from RM



# Application Workflow

→ Execution Sequence :


1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM
4. AM asks containers from RM
5. AM notifies NM to launch containers



# Application Workflow


→ Execution Sequence :

1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM
4. AM asks containers from RM
5. AM notifies NM to launch containers
6. Application code is executed in container




→ Execution Sequence :

1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM
4. AM asks containers from RM
5. AM notifies NM to launch containers
6. Application code is executed in container
7. Client contacts RM/AM to monitor application's status



→ Execution Sequence :

1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM
4. AM asks containers from RM
5. AM notifies NM to launch containers
6. Application code is executed in container
7. Client contacts RM/AM to monitor application's status
8. AM unregisters with RM



# Thank You

Slide content from  
Edureka and Simplilearn